

PromptRaise Blog

PRD - Regular Publishing + Internal Linking (Topical Authority Engine)

Status: Draft for approval · Date: 2026-08-12 · Owner: Zain Khan / PromptRaise · Prepared by: Zack (Hermes Agent)

PRD — PromptRaise Blog: Regular Publishing + Internal Linking (Topical Authority Engine)

Status: Draft for approval

Date: 2026-08-12

Owner: Zain Khan / PromptRaise

Prepared by: Zack (Hermes Agent)

1. Problem

- **Blog posts have no URLs.** `blog.promptraise.com` is a static Vite SPA where `src/App.tsx` swaps `ArticlePage` via React state — no router, no slugs, no crawlable URLs. Posts cannot be linked to, cannot be indexed, and cannot be cited by answer engines.

- **No regular publishing workflow.** Content is hardcoded in `src/pages/BlogPage.tsx` (POSTS array) and `src/pages/ArticlePage.tsx` (SECTIONS). Publishing a post requires editing code, rebuilding, and redeploying — nobody does that weekly.

- **Internal linking cannot happen.** The glossary (now 40 terms, live at `promptraise.com/academy/glossary`) has Phase 0 related-term chips, but Phase 1/2 (glossary ↔ blog posts, tool pages, service pages) is blocked because the blog has no stable URLs to link to.

- **Topical authority is capped.** AI engines build authority by crawling a dense, interlinked body of content. Without crawlable posts + a bidirectional link graph to the glossary, PromptRaise cannot compound topical authority around "Web3 AI visibility / GEO".

2. Goal

Give PromptRaise a **CMS-driven blog with real URLs** (`promptraise.com/blog`) that supports **regular publishing** (non-technical editor) and is wired into the **glossary internal-link graph** — turning the site into a topical-authority engine for AI visibility in Web3.

3. Non-goals (out of scope for V1)

- Migrating the old SPA blog content (we will copy the 6 existing posts in, not build a migration framework).
- Social sharing features, comments, author pages, newsletters.
- AI-writing automation inside the CMS (post-writing workflow stays manual; agent-assisted drafting is a separate project).
- Redirects from `blog.promptraise.com` URLs (they never had real URLs, so nothing to redirect).

4. Current state (verified)

---	---	---
Sanity `post` schema	`sanity/schemaTypes/postType.ts`	■ Exists (title, slug, excerpt, coverImage, categories, body blocks, author)
Sanity blog queries	`sanity/lib/queries.ts`	■ None — only pages/siteSettings
Blog routes	`app/blog/*`	■ None
Glossary data	`lib/glossary-terms.ts`	■ 40 terms + GLOSSARY_RELATED graph + relatedFor()/termAnchor()
Glossary pages	`app/academy/glossary`, `app/glossary`	■ Live, Phase 0 chips shipped
Old blog content	`/root/promptraise-blog/src`	6 posts + 1 featured, hardcoded

5. Proposed architecture

Option A (recommended): Sanity-backed blog on the main site

- New routes on `promptraise.com`:
 - `app/blog/page.tsx` — list page (all published posts, newest first)
 - `app/blog/[slug]/page.tsx` — post detail (SSG with `generateStaticParams` + ISR revalidate, matching the site's existing Sanity pattern)
- New Sanity queries in `sanity/lib/queries.ts`:
 - `getAllPosts()` (`published, status == "published", publishedAt <= now`)
 - `getPostBySlug(slug)` + draft-aware preview variant
 - Post body rendered from Sanity portable text blocks (reuse existing portable-text components if present, else add a minimal renderer).
- Blog list/detail get added to `sitemap.xml` and `llms.txt` automatically.
- Draft mode via existing `/api/draft` pattern (extend for posts).

Why Option A over keeping the Vite blog: one domain (authority concentration), one CMS (Sanity already powers the site), SSR (crawlable), ISR (instant publishes), and the glossary can deep-link to `promptraise.com/blog/[slug]` — all on the same origin.

6. Internal linking requirements (built into V1)

- **Glossary** → **blog**: add a `seeAlso: { label, href }[]` field to `GlossaryTerm`; populate for the terms that have a natural post (e.g. "AI visibility" → the "Why ChatGPT Ignores Your Protocol" post). Render as "Related resources" on the term card (both glossary pages).

- **Blog** → **glossary**: in post body copy, link the first use of any glossary-covered term to `promptraise.com/academy/glossary#term-<slug>`. Enforce via an editorial checklist (see §8) rather than auto-linking (auto-linkers produce bad anchor text).

- **Blog** ↔ **blog**: each post links to 2 related posts ("Keep reading") — done manually in the CMS body or via a related-posts query by category.

- Anchor text = exact term name (matches `DefinedTerm` name — sameAs consistency is what engines reward).

7. Publishing workflow (the "regular publishing" part)

- Editor writes post in Sanity Studio (`/studio`), status = Draft.
- Preview via Draft Mode (`/api/draft?slug=...`) — see it live.
- Editor flips status to Review → someone approves → flips to Published with `publishedAt`.

- ISR revalidation (`revalidate: 30` or `revalidateTag` on post publish) makes it live within ~30s; `sitemap` + `llms.txt` regenerate.

- Optional: IndexNow ping on publish (site already has this pattern for revalidate — extend to posts).

Cadence target: 1-2 posts/week. Content themes aligned to glossary categories (GEO, grounding, structured data, measurement, Web3-AI).

8. Editorial internal-link checklist (enforced per post)

- [] First use of each glossary term is linked to its `#term-<slug>` anchor
- [] At least 2 "related posts" linked inline or in a keep-reading block
- [] One link back to the glossary page itself
- [] One answer-capsule sentence at the top (quotable, 1 first-party number)
- [] Meta title ≤ 60 chars, meta description ≤ 160 chars

9. Acceptance criteria

- `promptraise.com/blog` lists published posts; each post has a crawlable URL `promptraise.com/blog/<slug>` returning 200 with SSR content.
- A post published in Sanity goes live without a code deploy, within ~1 min.
- `sitemap.xml` + `llms.txt` include blog URLs.
- At least 6 legacy posts migrated with proper internal links.

- Glossary "Related resources" links render on the live glossary.
- `curl` of any blog URL with JS disabled shows full article text

(AI Eyes test — crawlable without JS).

10. Risks / mitigations

- **Sanity dataset mismatch** (studio paused previously): re-verify datasets before build; if the studio is still paused, unblock it as part of this work.

- **Portable-text renderer missing**: check for existing components; if none, add a minimal block renderer (heading/para/code/image/link) — small.

- **Old posts have images on Unsplash**: reuse as-is, no migration needed.

- **Scope creep**: keep V1 to list + detail + links + publish flow. No categories/tags pages, no search, no authors.

11. Deliverables

- Blog list + detail routes (SSR/SSG, ISR)
- Sanity post queries + draft-mode wiring
- Blog URLs in sitemap + llms.txt
- Glossary `seeAlso` field + rendering (Phase 1 of the linking plan)
- 6 legacy posts migrated + internal links applied
- This PRD's acceptance criteria verified on staging

12. Approval gate

Per workflow: PRD approval (chat or PDF via email) → staging development

(`staging.zainkhan.me` staging env) → Zain alone publishes to production.